

# 1 Dynamic Programming

## 1.1 Knapsack Problem

Given a set of  $n$  items with integer sizes  $s_i$  and values  $v_i$ , and a knapsack of capacity  $C$ , find a set of items whose total size is less than  $C$ , but whose value is maximized. Items may be used more than once.

**Subproblem:**  $M(j)$ : the maximum value that can be packed into a size  $j$  knapsack

**Recurrence:**  $M(j) = \max\{M(j-1), \max_{1 \leq i \leq n} M(j-s_i) + v_i\}$

**Running time:**  $O(nC)$  time:  $C$  subproblems, each taking  $O(n)$  time.

## 1.2 0/1 Knapsack Problem

The knapsack problem, but it is forbidden to use more than one of each item.

**Subproblem:**  $M(i, j)$ : optimal value for filling *exactly* capacity  $j$  knapsack with a subset of items  $1, \dots, i$ .

**Recurrence:**  $M(i, j) = \max\{M(i-j, j), M(i-1, j-s_i) + v_i\}$

These represent not using the  $i$ th item, and using the  $i$ th item, respectively.

**Solution:**  $\max_j \{M(n, j)\}$

**Running time:**  $O(nC)$  time:  $nC$  subproblems that take  $O(1)$  time

## 1.3 Longest Path in a DAG

Given a DAG  $G = (V, E)$ , find the longest  $s \rightarrow t$  path.

**Subproblem:**  $d[u]$ : length of longest  $u \rightarrow t$  path

**Recurrence:**  $d[u] = \max_{(u,v) \in E} (w(u,v) + d[v])$

**Running Time:**  $\Theta(V + E)$ : Solve with a DFS on  $G$ .

## 1.4 Max Value Contiguous Subsequence

Given a sequence of  $n$  real numbers  $A_1, \dots, A_n$ , determine a contiguous subsequence  $A_i, \dots, A_j$  for which the  $\sum_{l=i}^j A_l$  is maximized.

**Subproblem:**  $M(j)$ : max window ending at  $j$

**Recurrence:**  $M(j) = \max\{M(j-1) + A_j, A_j\}$

**Solution:**  $\max_j M(j)$

**Running time:**  $O(n)$  time:  $n$  subproblems each take  $O(1)$  time

## 1.5 Making Change

Given coin denominations of values  $1 = v_1 < \dots < v_n$ , and an amount of money  $C$ , make  $C$  cents in change with as few coins as possible.

**Subproblem:**  $M(j)$ : minimum coins needed to make change for  $j$  cents

**Recurrence:**  $M(j) = \min_i \{M(j-v_i) + 1\}$

**Running time:**  $O(nC)$  time:  $C$  subproblems that take  $O(n)$  time

## 1.6 Longest Increasing Subsequence

Given a sequence of  $n$  real numbers  $A_1, \dots, A_n$ , find the longest subsequence (not necessarily contiguous) where values in the subsequence are strictly increasing.

**Subproblem:**  $L(j)$ : longest strictly increasing subsequence ending at position  $j$ .

**Recurrence:**  $L(j) = \max_{i < j, A_i < A_j} \{L(i)\} + 1$

**Solution:**  $\max_j L(j)$

**Running time:**  $O(n^2)$  time:  $n$  subproblems that each take  $n$  time

## 1.7 Edit Distance

Given strings  $A$  of length  $n$  and  $B$  of length  $m$ , transform  $A$  into  $B$  using a series of inserts, deletes, and replacements, of cost  $C_i$ ,  $C_d$ , and  $C_r$ , minimizing the cost.

**Subproblem:**  $T(i, j)$ : minimum cost to transform  $A[1:i]$  into  $B[1:j]$ .

**Recurrence:**

$$T(i, j) = \min \begin{cases} C_d + T(i-1, j) \\ T(i, j-1) + C_i \\ \begin{cases} T(i-1, j-1), & A[i] == B[j] \\ T(i-1, j-1), & \text{else} \end{cases} \end{cases}$$

**Running time**  $O(nm)$  time:  $nm$  subproblems that each take  $O(1)$  time

## 1.8 Counting Boolean Parenthizations

Given a boolean expression consisting of  $n$  Trues and Falses connected by either **and**, **or**, or **xor**, find the number of ways to parenthesize the expression such that it evaluates to **True**.

**Subproblems**  $T(i, j)$ : Number of ways to parenthesize booleans  $i$  through  $j$  such that the subexpression evaluates to **True**.

$F(i, j)$ : similar, but number of ways to make subexpression **False**

$S(i, j) = T(i, j) + F(i, j)$

**Recurrence**

$$T(i, j) = \sum_{i \leq k \leq j-1} \begin{cases} T(i, k)T(k+1, j) & \text{operator after } k \text{ is and} \\ S(i, k)S(k+1, j) - F(i, k)F(k+1, j) & \text{operator after } k \text{ is or} \\ T(i, k)F(k+1, j) + F(i, k)T(k+1, j) & \text{operator after } k \text{ is xor} \end{cases}$$

**Running time:**  $O(n^3)$ :  $n^2$  subproblems, each taking  $O(n)$  time

## 1.9 Subset Sum

Given a sequence of  $n$  numbers  $x_1, \dots, x_n$ , select the subset with the maximum total sum, subject to the constraint that you can't select two adjacent elements

**Subproblems**  $S(i)$ : maximum sum of the first  $i$  numbers

**Recurrence**  $S(i) = \max(S(i-2) + x_i, S(i-1))$

**Running time:**  $O(n)$  time:  $n$  subproblems, each taking  $O(1)$  time.

## 1.10 Reducable Problems

### 1.10.1 Box Stacking

Given a stack of 3-D boxes with height  $h_i$ , width  $w_i$ , and depth  $d_i$ , create the tallest stack of boxes possible, where box  $j$  can only be stacked on box  $i$  if  $w_i > w_j$  and  $d_i > d_j$ .

**Reduction:** Make a box for each possible rotation of the original set of boxes. Assume  $w_i \leq d_i$ . Sort boxes in order of decreasing base area. Problem is now Longest Increasing Subsequence

### 1.10.2 Building Bridges

Consider a map with a horizontal river. There are  $n$  cities on the southern bank with increasing x-coordinates  $A_1, \dots, A_n$ , and corresponding cities on the northern bank with x-coordinates (that are not necessarily increasing)  $B_1, \dots, B_n$ . Connect as many  $A_i, B_i$  pairs as possible without bridges crossing.

**Reduction:** Define a list with element  $i$  as the index of city  $i$  on the northern bank. Find the Longest Increasing Subsequence of that list.

### 1.10.3 Picking Up Pennies

Given a DAG  $G = (V, E)$  where some edges have pennies, find a path  $s \rightarrow t$  with the most pennies.

**Reduction:** Define  $w(u, v)$  = number of pennies on edge. Then Longest Path.

# 2 Asymptotics

$$\begin{aligned} f = O(g) &\iff \lim_{n \rightarrow \infty} f(n)/g(n) \neq \infty \\ f = \Omega(g) &\iff \lim_{n \rightarrow \infty} f(n)/g(n) \neq 0 \\ f = \Theta(g) &\iff \lim_{n \rightarrow \infty} f(n)/g(n) \notin \{0, \infty\} \\ f = \Theta(g) &\iff f = O(g) \text{ and } f = \Omega(g) \\ f = \Theta(g) &\iff g = \Theta(f) \\ f = O(g) &\iff g = \Theta(f) \end{aligned}$$

## 2.1 Common Functions

Ordered from slowest growing to fastest growing.

1	$\log \log n$
$\log n$	$(\log n)^c, c > 1$
$n^c, 0 < c < 1$	$n$
$n \log n = \log n!$	$n^2$
$n^3$	$n^c, c > 0$
$2^n$	$3^n$
$c^n, c > 1$	$n!$

## 3 Heaps

Invariant (for max-heap): if  $B$  is a child of  $A$  then  $A > B$ .

- **heapify**  $\Theta(\log n)$

Assumes that  $\text{LEFT}(i)$  and  $\text{RIGHT}(i)$  are heaps, but that  $A[i]$  may violate the heap invariant

```
def heapify(A, i):
    l = left(i)
    r = right(i)
    if l <= len(A) and A[l] > A[i]:
        largest = l
    else:
        largest = i

    if r <= len(A) and A[r] > A[largest]:
        largest = r

    if largest != i:
        A[i], A[largest] = A[largest], A[i]
        heapify(A, largest)
```

- **extract-max**  $\Theta(\log n)$

```
def extract_max(A):
    max = A[0]
    A[0] = A.pop(-1)
    heapify(A, 0)
    return max
```

- **increase-key**  $\Theta(\log n)$

```
def increase_key(A, i, key):
    A[i] = key
    while i > 0 and A[parent(i)] < A[i]:
        A[i], A[parent(i)] = A[parent(i)], A[i]
        i = parent(i)
```

- **insert** Add  $-\infty$  to the end of the heap, then **increase-key** to the actual value

## 4 Sorting

Algorithm	Best	Worst	Stable	In-place
Insertion Sort	$n$	$n^2$	Y	Y
Merge Sort	$n \lg n$	$n \lg n$	Y	N
Heap Sort	$n$	$n \lg n$	N	Y
Counting Sort	$n$	$n$	Y	N
Quicksort	$n \lg n$	$n^2$	some versions	N

- **Insertion Sort**

Very good on sorted sets; online; good for small sets. Atrocious for large sets.

```
INSERTION-SORT(A)
1 for j <- 2 to length[A]
2   do key <- A[j]
3   Insert A[j] into the sorted sequence A[1... j - 1].
4   i <- j - 1
5   while i > 0 and A[i] > key
6     do A[i + 1] <- A[i]
7     i <- i - 1
8   A[i + 1] <- key
```

- **Merge Sort**

Wikipedia: "Although heapsort has the same time bounds as merge sort, it requires only  $\Theta(1)$  auxiliary space instead of merge sort's  $\Theta(n)$ , and is often faster in practical implementations. Quicksort, however, is considered by many to be the fastest general-purpose sort algorithm. On the plus side, merge sort is a stable sort, parallelizes better, and is more efficient at handling slow-to-access sequential media. Merge sort is often the best choice for sorting a linked list: in this situation it is relatively easy to implement a merge sort in such a way that it requires only  $\Theta(1)$  extra space, and the slow random-access performance of a linked list makes some other algorithms (such as quicksort) perform poorly, and others (such as heapsort) completely impossible."

```
MERGE-SORT(A, p, r)
1 if p < r
2   then q <- (floor((p + r)/2))
3   MERGE-SORT(A, p, q)
4   MERGE-SORT(A, q + 1, r)
5   MERGE(A, p, q, r)
```

```
function merge(left, right)
    var list result
    while length(left) > 0 and length(right) > 0
        if first(left) <= first(right)
            append first(left) to result
            left = rest(left)
        else
            append first(right) to result
            right = rest(right)
    end while
    while length(left) > 0
        append left to result
    while length(right) > 0
        append right to result
    return result
```

- **Heap Sort**

Algorithm in a nutshell: 1. Build tree. 2. Grab root. 3. Max-heapify. 4. Go to 2. Not stable, but the  $n \lg n$  worst case is nice. Usually slower than quicksort.

```
HEAPSORT(A)
1 BUILD-MAX-HEAP(A)
2 for i <- length[A] downto 2
3   do exchange A[1] <-> A[i]
4   heap-size[A] <- heap-size[A] - 1
5   MAX-HEAPIFY(A, 1)
```

- **Counting Sort**

```
COUNTING-SORT(A, B, k)
1 for i <- 0 to k
2   do C[i] <- 0
3 for j <- 1 to length[A]
4   do C[A[j]] <- C[A[j]] + 1
5 C[i] now contains the number of elements equal \
  to i.
6 for i <- 1 to k
7   do C[i] <- C[i] + C[i - 1]
8 C[i] now contains the number of elements less \
  than or equal to i.
9 for j <- length[A] downto 1
10  do B[C[A[j]]] <- A[j]
11  C[A[j]] <- C[A[j]] - 1
```

- **Quicksort**

```
QUICKSORT(A, p, r)
1 if p < r
2   then q <- PARTITION(A, p, r)
3   QUICKSORT(A, p, q - 1)
4   QUICKSORT(A, q + 1, r)
```

```
PARTITION(A, p, r)
1 x <- A[r]
2 i <- p - 1
3 for j <- p to r - 1
4   do if A[j] <= x
5     then i <- i + 1
6     exchange A[i] <-> A[j]
7 exchange A[i + 1] <-> A[r]
8 return i + 1
```

- **Radix sort**

```
RADIX-SORT(A, d)
1 for i <- 1 to d
2   do use a stable sort to sort array A on digit i
```

Given  $n$   $d$ -digit numbers in which each digit can take on up to  $k$  possible values, RADIX-SORT correctly sorts these numbers in  $\Theta(d(n+k))$  time. Stable; not in-place.

- **Bucket Sort**

Assumes that the input is generated by a random process that distributes elements uniformly over  $[0,1)$

```

BUCKET-SORT(A)
1  n <- length[A]
2  for i <- 1 to n
3      do insert A[i] into list B[floor(n A[i])]
4  for i <- 0 to n - 1
5      do sort list B[i] with insertion sort
6  concatenate the lists B[0], B[1], ..., B[n - 1] together in order

```

## 5 Search

- $d[u]$  is time  $u$  is discovered (added to queue)  
 $f[u]$  is time that  $u$  is finished (removed from queue)
- Tree edges are edges in the depth-first forest  $G_p$ . Edge  $(u, v)$  is a tree edge if  $v$  was first discovered by exploring edge  $(u, v)$ . Edge  $(u, v)$  is a tree or forward edge iff  $d[u] < d[v] < f[v] < f[u]$
- Back edges are those edges  $(u, v)$  connecting a vertex  $u$  to an ancestor  $v$  in a depth-first tree. Self-loops, which may occur in directed graphs, are considered to be back edges. Edge  $(u, v)$  is a back edge if and only if  $d[v] < d[u] < f[u] < f[v]$
- Forward edges are those nontree edges  $(u, v)$  connecting a vertex  $u$  to a descendant  $v$  in a depth-first tree.
- Cross edges are all other edges. They can go between vertices in the same depth-first tree, as long as one vertex is not an ancestor of the other, or they can go between vertices in different depth-first trees. Edge  $(u, v)$  is a cross edge if and only if  $d[v] < f[v] < d[u] < f[u]$ .
- Diameter is the maximum-weight shortest path.
- A connected graph has a minimum of  $V - 1$  edges and a maximum of  $V^2/2$  if undirected, twice that if not. A connected graph has  $V^2/2$  if undirected, etc. A sparse graph has fewer than  $O(V)$  edges.

### 5.1 Topological Sort

A topological sort of a dag  $G = (V, E)$  is a linear ordering of all its vertices such that if  $G$  contains an edge  $(u, v)$ , then  $u$  appears before  $v$  in the ordering. (If the graph is not acyclic, then no linear ordering is possible.) Topological orderings are not necessarily unique

```

TOPOLOGICAL-SORT(G)
1  call DFS(G) to compute finishing times f[v] for each vertex v
2  as each vertex is finished, insert it onto the front of a linked list
3  return the linked list of vertices

```

We can perform a topological sort in time  $\Theta(V + E)$ , since depth-first search takes  $\Theta(V + E)$  time and it takes  $O(1)$  time to insert each of the  $|V|$  vertices onto the front of the linked list.

### 5.2 Shortest Path

#### • Shortest path - general

In computer science, a problem is said to have optimal substructure if an optimal solution can be constructed efficiently from optimal solutions to its subproblems. (In this context: shortest paths between two points have other shortest paths in them)

If there is a negative-weight cycle reachable from  $s$ , however, shortest-path weights are not well defined. No path from  $s$  to a vertex on the cycle can be a shortest path—a lesser-weight path can always be found that follows the proposed “shortest” path and then traverses the negative-weight cycle. If there is a negative-weight cycle on some path from  $s$  to  $v$ , we define  $\delta(s, v) = -\infty$

Triangle inequality: Cutting out an intermediate stop never increases the cost. We formalize this notion by saying that the cost function  $c$  satisfies the triangle inequality if for all vertices  $u, v, w \in V$ ,  $c(u, w) \leq c(u, v) + c(v, w)$

Predecessor subgraph: all  $\pi[i]$

#### • Relaxation

The process of relaxing[1] an edge  $(u, v)$  consists of testing whether we can improve the shortest path to  $v$  found so far by going through  $u$  and, if so, updating  $d[v]$  and  $\pi[v]$ . A relaxation step may decrease the value of the shortest-path estimate  $d[v]$  and update  $v$ 's predecessor field  $\pi[v]$ . The following code performs a relaxation step on edge  $(u, v)$ .

```

RELAX(u, v, w)
1  if d[v] > d[u] + w(u, v)
2      then d[v] <- d[u] + w(u, v)
3      pi[v] <- u

```

#### • Bellman-Ford

The Bellman-Ford algorithm solves the single-source shortest-paths problem in the general case in which edge weights may be negative. Given a weighted, directed graph  $G = (V, E)$  with source  $s$  and weight function  $w : E \rightarrow R$ , the Bellman-Ford algorithm returns a boolean value indicating whether or not there is a negative-weight cycle that is reachable from the source. If there is such a cycle, the algorithm indicates that no solution exists. If there is no such cycle, the algorithm produces the shortest paths and their weights. If the graph does contain a cycle of negative weights, Bellman-Ford can only detect this; Bellman-Ford cannot find the shortest path that does not repeat any vertex in such a graph. This problem is at least as hard as the NP-complete longest path problem.

The algorithm uses relaxation, progressively decreasing an estimate  $d[v]$  on the weight of a shortest path from the source  $s$  to each vertex  $v \in V$  until it achieves the actual shortest-path weight  $\delta(s, v)$ . The algorithm returns TRUE if and only if the graph contains no negative-weight cycles that are reachable from the source.

```

BELLMAN-FORD(G, w, s)
1  INITIALIZE-SINGLE-SOURCE(G, s)
2  for i <- 1 to |V[G]| - 1
3      do for each edge (u, v) in E[G]
4          do RELAX(u, v, w)
5  for each edge (u, v) in E[G]
6      do if d[v] > d[u] + w(u, v)
7          then return FALSE
8  return TRUE

```

The Bellman-Ford algorithm runs in time  $O(VE)$ , since the initialization in line 1 takes  $\Theta(V)$  time, each of the  $|V| - 1$  passes over the edges in lines 2-4 takes  $\Theta(E)$  time, and the for loop of lines 5-7 takes  $O(E)$  time.

#### • Dijkstra

Dijkstra's algorithm solves the single-source shortest-paths problem on a weighted, directed graph  $G = (V, E)$  for the case in which all edge weights are nonnegative. In this section, therefore, we assume that  $w(u, v) \geq 0$  for each edge  $(u, v) \in E$ . As we shall see, with a good implementation, the running time of Dijkstra's algorithm is lower than that of the Bellman-Ford algorithm.

Dijkstra's algorithm maintains a set  $S$  of vertices whose final shortest-path weights from the source  $s$  have already been determined. The algorithm repeatedly selects the vertex  $u \in V - S$  with the minimum shortest-path estimate, adds  $u$  to  $S$ , and relaxes all edges leaving  $u$ . In the following implementation, we use a min-priority queue  $Q$  of vertices, keyed by their  $d$  values.

Runtime:

- If we simply store  $d[v]$  in the  $v$ th entry of an array. Each INSERT and DECREASE-KEY operation takes  $O(1)$  time, and each EXTRACT-MIN operation takes  $O(V)$  time (since we have to search through the entire array), for a total time of  $O(V^2 + E) = O(V^2)$ .
- If the graph is sufficiently sparse in particular,  $E = o(V^2/\lg V)$  it is practical to implement the min-priority queue with a binary min-heap. Each EXTRACT-MIN operation then takes time  $O(\lg V)$ . As before, there are  $|V|$  such operations. The time to build the binary min-heap is  $O(V)$ . Each DECREASE-KEY operation takes time  $O(\lg V)$ , and there are still at most  $|E|$  such operations. The total running time is therefore  $O((V + E)\lg V)$ , which is  $O(E\lg V)$  if all vertices are reachable from the source. This running time is an improvement over the straightforward  $O(V^2)$ -time implementation if  $E = o(V^2/\lg V)$ .
- We can in fact achieve a running time of  $O(V \lg V + E)$  by implementing the min-priority queue with a Fibonacci heap (see Chapter 20). The amortized cost of each of the  $|V|$  EXTRACT-MIN operations is  $O(\lg V)$ , and each DECREASE-KEY call, of which there are at most  $|E|$ , takes only  $O(1)$  amortized time.

```

DIJKSTRA(G, w, s)
1  INITIALIZE-SINGLE-SOURCE(G, s)
2  S <- null
3  Q <- V[G]
4  while Q != null
5      do u <- EXTRACT-MIN(Q)
6         S <- S ∪ {u}
7         for each vertex v in Adj[u]
8             do RELAX(u, v, w)

```




Figure 24.1: Negative edge weights in a directed graph. Shown within each vertex is its shortest-path weight from source  $s$ . Because vertices  $e$  and  $f$  form a negative-weight cycle reachable from  $s$ , they have shortest-path weights of  $-\infty$ . Because vertex  $g$  is reachable from a vertex whose shortest-path weight is  $-\infty$ , it, too, has a shortest-path weight of  $-\infty$ . Vertices such as  $h$ ,  $i$ , and  $j$  are not reachable from  $s$ , and so their shortest-path weights are  $-\infty$ , even though they lie on a negative-weight cycle.




Figure 24.4: The execution of the Bellman-Ford algorithm. The source is vertex  $s$ . The  $d$  values are shown within the vertices, and shaded edges indicate predecessor values: if edge  $(u, v)$  is shaded, then  $\pi[v] = u$ . In this particular example, each pass relaxes the edges in the order  $(t, x)$ ,  $(t, y)$ ,  $(t, z)$ ,  $(x, t)$ ,  $(y, x)$ ,  $(y, z)$ ,  $(z, x)$ ,  $(z, s)$ ,  $(s, t)$ ,  $(s, y)$ . (a) The situation just before the first pass over the edges. (b)-(e) The situation after each successive pass over the edges. The  $d$  and  $\pi$  values in part (e) are the final values. The Bellman-Ford algorithm returns TRUE in this example.




Figure 24.5: The execution of the algorithm for shortest paths in a directed acyclic graph. The vertices are topologically sorted from left to right. The source vertex is  $s$ . The  $d$  values are shown within the vertices, and shaded edges indicate the values. (a) The situation before the first iteration of the for loop of lines 35. (b)-(g) The situation after each iteration of the for loop of lines 35. The newly blackened vertex in each iteration was used as  $u$  in that iteration. The values shown in part (g) are the final values.




Figure 24.3: Relaxation of an edge  $(u, v)$  with weight  $w(u, v) = 2$ . The shortest-path estimate of each vertex is shown within the vertex. (a) Because  $d[v] > d[u] + w(u, v)$  prior to relaxation, the value of  $d[v]$  decreases. (b) Here,  $d[v] \leq d[u] + w(u, v)$  before the relaxation step, and so  $d[v]$  is unchanged by relaxation.




Figure 24.6: The execution of Dijkstra's algorithm. The source  $s$  is the leftmost vertex. The shortest-path estimates are shown within the vertices, and shaded edges indicate predecessor values. Black vertices are in the set  $S$ , and white vertices are in the min-priority queue  $Q = V - S$ . (a) The situation just before the first iteration of the while loop of lines 48. The shaded vertex has the minimum  $d$  value and is chosen as vertex  $u$  in line 5. (b)-(f) The situation after each successive iteration of the while loop. The shaded vertex in each part is chosen as vertex  $u$  in line 5 of the next iteration. The  $d$  and  $\pi$  values shown in part (f) are the final values.

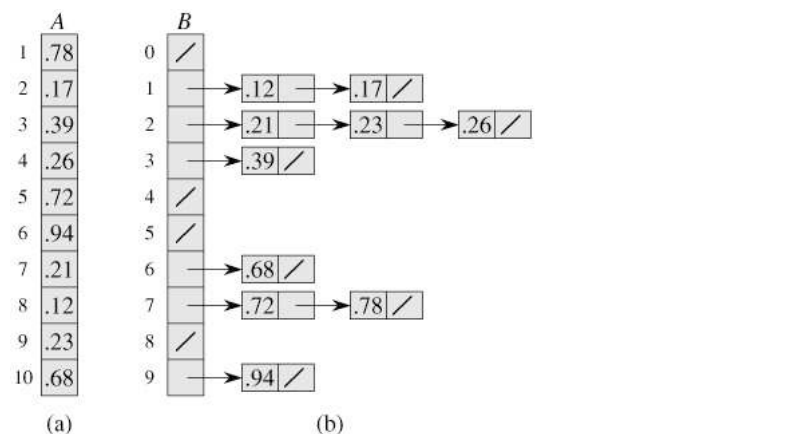




Figure 8.4: The operation of BUCKET-SORT. (a) The input array  $A[1..10]$ . (b) The array  $B[0..9]$  of sorted lists (buckets) after line 5 of the algorithm. Bucket  $i$  holds values in the half-open interval  $[i/10, (i+1)/10)$ . The sorted output consists of a concatenation in order of the lists  $B[0], B[1], \dots, B[9]$ .

TREES/BSTs



- A. INSERT(78)  $\Theta(h)$
- B. DELETE(50) — swap with 61 and delete  $\Theta(h)$
- C. SUCCESSOR  $\Theta(h)$
- D. IN-ORDER WALK  $\Theta(n)$
- AND: CONSTRUCTING A BST FROM A SORTED LIST... (also  $\Theta(n)$ )

HEAPS



NEARLY-COMPLETE BINARY TREE

MAX-HEAP... OF COURSE, MIN-HEAPS EXIST, TOO.

- A. EXTRACT-MIN\*  $\Theta(\lg n)$
- B. HEAPIFY  $\Theta(\lg n)$
- C. BUILD-HEAP  $\Theta(n)$
- D. PRIORITY QUEUE — DECREASE-KEY\*


ARRAY REPRESENTATION:

90	74	82	61	2	...
0	1	2	3	4	...

PARENT[j] =  $\lfloor \frac{j-1}{2} \rfloor$

CHILD[i]:  
LEFT[i] =  $2i + 1$       RIGHT[i] =  $2i + 2$


AVL TREES: BALANCED BSTs




$1 \geq |h(\text{LEFT}) - h(\text{RIGHT})|$

REBALANCING ONLY REQUIRES ROTATIONS



EASY CASE:




HARD CASE:




THE ABOVE WON'T WORK... WE NEED TWO ROTATIONS



THEN IT BECOMES THE EASY CASE...



# HASHING / HASH TABLES



A. COLLISION RESOLUTION

B. INSERT

C. DELETE

D. LOOKUP

E. ROLLING HASHES

EXPECTED  $\Theta(1)$   
WORST CASE  $\Theta(n)$

LOAD FACTOR  $\alpha = \frac{n}{m}$   
(n ELTS. INSERTED)

## AMORTIZED ANALYSIS

HASH TABLE (IGNORE DELETIONS) — WANT TO KEEP  $\alpha < \frac{4}{5}$

ON INSERT, IF  $\alpha \geq \frac{4}{5}$ , ALLOCATE A NEW TABLE OF  $2m$  SIZE AND REHASH EVERYTHING  
 $\rightarrow O(m+n) = O(n)$  TIME TO RESIZE  
 $O(1)$  OTHERWISE

IN A SERIES OF  $n$  INSERTIONS:

- suppose resizes occur at  $k, 2k, 4k, \dots$

- then cost of insertions is  $O(n + (k + 2k + \dots + n)) = O(n + (\frac{n}{2^1} + \frac{n}{2^{i-1}} + \dots + n))$   
 $= O(n + (2n))$   
 $= O(3n) = O(n)$

Thus each insert is  $\frac{O(n)}{n} = O(1)$  time